

1 **A Experiments compute resources**

- 2 All experiments were conducted on a server running Ubuntu 22.04.4 LTS. The system is equipped
 3 with an Intel Xeon Gold 6133 processor featuring 80 logical cores at 3.0GHz, and 352 GB of RAM.
 4 The machine is configured with 8 NVIDIA GeForce RTX 4080 GPUs with 16GB of video memory
 5 each.

6 **B Detail of Mamba**

- 7 The following provides a detailed explanation of the internal computations within the Mamba module,
 8 beginning with the specification of the corresponding dimensional representations.

$$\begin{aligned} B &:= \text{Batch size}, \\ S &:= \text{Sequence length}, \\ D_m &:= \text{dimension of model}, \\ D_{in} &:= \text{dimension of inner model} \\ &= \text{expand} \times D_m, \\ N &:= \text{dimension of hidden state}, \\ Num_h &:= \text{number of heads}, \\ H_d &:= \text{dimension of heads}, \end{aligned}$$

- 9 The computations within the Mamba module can be divided into three components: pre-SSM, SSM,
 10 and post-SSM. The following presents a detailed derivation for each submodule.

$$\text{Input } u \in \mathbb{R}^{(B, S, D_m)},$$

Pre-SSM

$$\begin{aligned} zxBCdt &= \mathbf{Linear}_{\text{proj}}(u) \in \mathbb{R}^{(B, S, 2D_{in}+2N+Num_h)}, \\ (z, xBC, dt) &= zxBCdt \in \mathbb{R}^{(B, S, (D_{in}, D_{in}+2N, Num_h))}, \\ x_c B_c C_c &= \mathbf{Conv1d}(xBC) \in \mathbb{R}^{(B, S, D_{in}+2N)}, \\ x_{cf} B_{cf} C_{cf} &= \mathbf{SiLU}(xBC) \in \mathbb{R}^{(B, S, D_{in}+2N)}, \\ (x_{cf}, B_{cf}, C_{cf}) &= x_{cf} B_{cf} C_{cf} \in \mathbb{R}^{(B, S, (D_{in}, N, N))}, \\ A_{init} &\sim \mathbf{Uniform}(A_{\min}, A_{\max}) \in \mathbb{R}^{Num_h}, \\ A_{\log} &= \log(A_{init}) \in \mathbb{R}^{Num_h}, \\ A &= -\exp(A_{\log}) \in \mathbb{R}^{Num_h}, \\ dtb_{init} &= \exp(\mathbf{rand}(Num_h) \cdot (\log(dt_max) - \log(dt_min)) + \log(dt_min)) \in \mathbb{R}^{Num_h}, \\ dt_{bias} &= dtb_{init} + \log(-\exp(-dtb_{init}) + 1) \in \mathbb{R}^{Num_h}, \\ dt_{fb} &= \mathbf{softplus}(dt + dt_{bias}) \in \mathbb{R}^{(B, S, Num_h)}, \\ \tilde{A} &= A \circ dt_{fb} \in \mathbb{R}^{(B, S, Num_h)}, \\ \tilde{x} &= x_{cf}.\mathbf{reshape}(B, S, Num_h, H_d) \circ dt_{fb}.\mathbf{reshape}(B, S, Num_h, 1) \in \mathbb{R}^{(B, S, Num_h, H_d)}, \\ \tilde{B} &= B_{cf}.\mathbf{reshape}(B, S, 1, N) \in \mathbb{R}^{(B, S, 1, N)}, \\ \tilde{C} &= C_{cf}.\mathbf{reshape}(B, S, 1, N) \in \mathbb{R}^{(B, S, 1, N)}. \end{aligned}$$

- 11 Here, **Linear**_{proj} denotes a linear transformation, **Conv1d** refers to a one-dimensional convolution,
 12 and **SiLU** represents the SiLU activation function. **Uniform** indicates sampling from a uniform
 13 distribution, while **softplus** denotes the Softplus activation function. **rand** refers to drawing a
 14 specified number of random values. The \circ symbol denotes element-wise multiplication, and **reshape**
 15 indicates a dimensional transformation operation.

SSM

$$\begin{aligned}
\hat{A} &= \text{Mask}_1 \circ \text{repeat}(\tilde{A}) \in \mathbb{R}^{(B, \text{Num}_h, S, S)}, \\
A_{cumsum} &= \text{cumsum}(\hat{A}) \in \mathbb{R}^{(B, \text{Num}_h, S, S)}, \\
L &= \exp(\text{Mask}_2 \circ A_{cumsum}) \in \mathbb{R}^{(B, \text{Num}_h, S, S)}, \\
P &= \text{einsum}("BSHN, BSHN \rightarrow BHSS", \tilde{C}, \tilde{B}) \in \mathbb{R}^{(B, 1, S, S)}, \\
M &= \text{einsum}("BHSS, BHSS \rightarrow BHSS", L, P) \in \mathbb{R}^{(B, \text{Num}_h, S, S)}, \\
y &= \text{einsum}("BHSS, BSHP \rightarrow BSHP", M, \tilde{x}) \in \mathbb{R}^{(B, S, \text{Num}_h, H_d)}.
\end{aligned}$$

16 The **Mask**₁ operation is used to zero out the diagonal and upper-triangular elements of a matrix,
17 retaining only the elements below the diagonal. The **Mask**₂ sets the elements above the diagonal
18 to negative infinity. The **cumsum** operation performs a cumulative sum along the first of the last
19 two dimensions (i.e., across rows) in a matrix of shape (S, S) , summing from top to bottom. The
20 **einsum** operation denotes the Einstein summation convention, used for concise and flexible tensor
21 contractions.

Post-SSM

$$\begin{aligned}
D &= \mathbf{1}_{D_{in}} \in \mathbb{R}^{D_{in}}, \\
\dot{y} &= y + x_{cf}.\text{reshape}(B, S, \text{Num}_h, H_d) \times D.\text{reshape}(D_{in}, 1) \in \mathbb{R}^{(B, S, \text{Num}_h, H_d)}, \\
\ddot{y} &= \dot{y}.\text{reshape}(B, S, D_{in}) \in \mathbb{R}^{(B, S, D_{in})}, \\
y_z &= \ddot{y} \cdot \text{SiLU}(z) \in \mathbb{R}^{(B, S, D_{in})}, \\
y_{norm} &= y_z \cdot \frac{1}{\sqrt{\text{mean}(y_z^2, \text{axis} = -1)} + \epsilon} \cdot \mathbf{w} \in \mathbb{R}^{(B, S, D_{in})}, \\
y_{out} &= \text{Linear}_{\text{proj}}(y_{norm}) \in \mathbb{R}^{(B, S, D_m)}.
\end{aligned}$$

22 Here, **1** denotes a vector in which all elements are equal to 1.

$$\text{Output } y_{out} \in \mathbb{R}^{(B, S, D_m)}.$$

23 C Data setup

24 C.1 Composite function task

25 **Standard** The total dataset comprises 300,000 samples, and each sequence has a fixed length of
26 8. Each anchor pair in the training set accounts for 5.6% of the total data, while each anchor pair
27 in the test set constitutes 0.6%. The mapping between anchors and their associated functions is as
28 follows: anchor 1 corresponds to a shift of +5, anchor 2 to +1, anchor 3 to -2, and anchor 4 to -8.
29 During training, all 15 possible anchor pairs are included except for pair 43. Among these, all pairs
30 except 34 are derived from the composition of their corresponding single-anchor functions. Notably,
31 the function for pair 34 is manually set to -6, deviating from the correct compositional result of
32 -10. The test set contains both symmetric and compositional instances of pair 43. The distinction
33 between training and test data is governed by the position of the key token. For sequences of length 8,
34 each position is associated with a congruence class modulo 8. In the training set, the key token is
35 prohibited from appearing in the position whose modulo-8 class matches its own. For example, key
36 33 cannot appear in the first position, as it belongs to the congruence class 1 mod 8, which is aligned
37 with the first index. Conversely, in the test set, each key token is required to occupy the position
38 that corresponds exactly to its modulo-8 congruence class. This design ensures that the model has
39 access to the full semantic range of tokens during training while preventing it from relying solely on
40 positional memorization to generalize to the test set.

41 **Full symmetry** The total number of datasets is 300,000, and each sequence has a fixed length
42 of 8. Each anchor pair in the training set accounts for 4.5% of the total, while each anchor in the
43 test set accounts for 0.5% of the total. To ensure that the only possible solutions are symmetric,

44 we simultaneously utilize two sets of correspondences between anchors and functions. This setup
 45 prevents the model from simply solving for individual anchor functions and instead forces it to derive
 46 symmetric solutions by understanding the symmetry of anchor pairs. There are a total of five anchors:
 47 0, 1, 2, 3, and 4. The 0 anchor is added to balance the data volume between the two function sets, thus
 48 avoiding model bias caused by disparities in data quantity. The first set of correspondences between
 49 anchors and functions is as follows: 0 corresponds to +2, 1 corresponds to +5, 2 corresponds to +1,
 50 3 corresponds to -2, and 4 corresponds to -8. The second set of correspondences is: 0 corresponds
 51 to -9, 1 corresponds to +6, 2 corresponds to -7, and 3 corresponds to +3. All anchor pair functions
 52 are composed of either the first or the second set of correspondences. To facilitate the observation of
 53 symmetric solutions, anchor pairs 00, 11, 22, 33, and 44 are excluded from this task. Thus, there are
 54 a total of 20 anchor pairs in this task. The anchor pairs following the first set of correspondences are
 55 01, 02, 10, 20, 14, 41, 23, 32, 34, 43, a total of 10 pairs. Among them, there are 4 1s, 4 ones, 4 twos,
 56 4 threes and 4 fours. This design ensures that each anchor has the same amount of data under the
 57 two corresponding methods, and the two corresponding methods also have the same amount of data.
 58 The anchor pairs following the second set of correspondences are 03, 30, 04, 40, 12, 21, 13, 31, 24,
 59 42. Only 43 is not included in the training set. The purpose of this task is to observe whether the
 60 model can discover the symmetry among all anchor pairs and thus output symmetric solutions. The
 61 distinction between the training set and the test set in the dataset is based on the position of the key.
 62 For sequences of length 8, each position corresponds to a congruence class modulo 8. In the training
 63 set, for each key, the key does not fall within the congruence class corresponding to its position. For
 64 example, the key 33 cannot appear in the first position of the sequence because it would then be in the
 65 congruence class modulo 8 of 1, which corresponds to that position. In contrast, in the test set, each
 66 key must exactly fall within the congruence class corresponding to its position. This setup allows the
 67 model to learn the meaning of all tokens while preventing it from relying solely on memorization to
 68 generalize to the test set.

69 C.2 Inverse sequence matching task

70 The total dataset consists of 100,000 samples. The sequence length varies with the number of layers
 71 in the Mamba model. Each sequence is generated from a generation set consisting of three distinct
 72 elements. It contains five different permutations of the elements in the generation set, where each
 73 permutation is followed by a random token that does not belong to the generation set. Each such
 74 permutation is referred to as a key sequence.

75 One of these five sequences is randomly selected and reversed to form the query sequence, which is
 76 appended to the end of the original sequence. Between the key sequences and the query sequence,
 77 a number of randomly generated tokens—equal to the convolutional receptive field of Mamba are
 78 inserted to prevent the model from retrieving information through convolution. Therefore, the total
 79 sequence length can be formally expressed as follows:

$$S = 5 \times (3 + 1) + 3 \times \text{Num}_{layer} + 3.$$

80 Here, Num_{layer} denotes the number of layers in the Mamba model. For example, in the case of
 81 a two-layer Mamba model, the sequence length is 29. To ensure a fair comparison, Mamba and
 82 Transformer models with the same number of layers are trained on identical sequence configurations.

83 The training set constitutes 80% of the total dataset, while the test set and the out-of-distribution
 84 (OOD) set each account for 10%. In the training set, the three elements in the generating set do not
 85 all belong to the same congruence class modulo 3. In contrast, in the test set, all three elements in the
 86 generating set belong to the same congruence class modulo 3. For both the training and test sets, all
 87 numerical values in the sequences are drawn from the range 20 to 100. In the OOD set, however, all
 88 sequence elements are drawn from the range 101 to 200, and the generation set is not subject to any
 89 congruence constraints.

90 D Training setup

91 Unless otherwise specified, all tasks in this work adopt the following training parameter settings. The
 92 learning rate is initially set to 1e-5 at the start of training, warmed up to 25 times its initial value
 93 within 10 epochs, and then decreases to 1e-5 via cosine decay at 200 epochs. The training optimizer
 94 is AdamW with parameters set as $\beta_1 = 0.9$, $\beta_2 = 0.999$, $eps = 1e - 8$, and weight decay=1e-2.

95 Meanwhile, gradient clipping is applied with a maximum norm of 1. For composite function tasks,
96 the batch size is set to 2048, while for inverse sequence matching tasks, the batch size is 1024. The
97 loss function used for training is the cross - entropy loss function, which is calculated only for the
98 last token of the model output sequence.

99 **E Experiment detail and result**

100 **E.1 Composite function task**

101 **E.1.1 Phase diagram of Mamba on the composite function task**

102 **Mamba configurations** The dimension of model is set to 32, the dimension of hidden state is
103 set to the default value of 128, and the expansion factor is set to the default value of 2, and the
104 activation function employed throughout the model is the Sigmoid Linear Unit (SiLU). To facilitate
105 clear observation, all experiments are conducted using a single head. The convolution kernel length
106 is set to its default value of 4. The values reported at different positions represent the mean results
107 obtained using three fixed random seeds. Across these positions, the only variations are in the model’s
108 depth or initialization.

109 **E.1.2 Experiments with all-one convolution**

110 **Mamba configurations** To investigate whether the asymmetry of convolution is responsible for
111 Mamba’s difficulty in learning the symmetric solution, we select a model configuration under
112 which Mamba successfully learns the compositional solution but fails to learn the symmetric one.
113 Specifically, we use a five-layer Mamba model with small initialization($\gamma = 1$), while all other
114 settings are consistent with those used in the phase diagram experiments.

115 **E.1.3 Experiments with positional encoding**

116 **Mamba configurations** To examine whether Mamba’s bias for asymmetry in the composite
117 function task arises from its lack of explicit positional encoding, thus encouraging reliance on
118 convolution. We select a configuration under which Mamba struggles to learn both the composite and
119 symmetric solutions. Specifically, we use a two-layer Mamba model with standard initialization($\gamma =$
120 0.5). All other settings are kept consistent with those used in the phase diagram experiments.

121 **E.1.4 Experiment under full symmetry**

122 **Mamba configurations** To validate Mamba’s difficulty in solving the composite function task
123 under fully symmetric settings, we set dimension of model to 128, while keeping all other settings
124 consistent with those used in the phase diagram experiments. The number of layers is varied across 2,
125 3, 4, 5, and 6, and standard initialization($\gamma = 0.5$) is applied.

126 **Extended results** For each configuration, experiments are conducted using three fixed random
127 seeds. The results are shown in the Fig. 1 and Fig. 2. As can be observed, under all configurations,
128 Mamba consistently struggles to solve the composite function task in the fully symmetric setting.

129 **E.2 Inverse sequence matching task**

130 **Mamba configurations** The dimension of model is set to 128, the dimension of hidden state is set
131 to the default value of 128, and the expansion factor is set to the default value of 2. The activation
132 function used is SiLU. To clearly isolate the structural differences between Transformer and Mamba
133 and ensure a fair comparison, the Transformer model also adopts SiLU as its activation function. The
134 convolution kernel length is set to its default value of 4, and the number of heads is fixed at 1.

135 **Transformer configurations** The dimension of model is set to 128, consistent with the Mamba
136 configuration. To ensure a fair comparison, the Transformer also adopts a $2 \times$ dimensional expansion
137 when generating value vectors, mirroring Mamba’s setup. Specifically, the dimensions of the query
138 and key vectors are set to 128, while the value vectors have a dimension of 256. Additionally, to
139 ensure a fair comparison and maintain a comparable number of parameters between the two models,




Figure 1: Accuracy of Mamba under fully symmetric setting. The horizontal axis represents the random seed, while the vertical axis corresponds to the number of layers in the Mamba model.

140 the Transformer’s feedforward network (FNN) uses a hidden dimension of 128, and, like Mamba, it
 141 is configured with only a single attention head throughout.

142 **Extended results** All experiments are conducted using the same set of random seeds to ensure
 143 fairness in comparison.

144 To evaluate Mamba’s difficulty in solving the inverse sequence matching task, we vary the model
 145 depth across 2, 3, 4, and 5 layers, and consider different initialization strategies ranging from standard
 146 to small initialization($\gamma = 0.5, 0.6, 0.7, 0.8, 0.9, 1.0$). The detailed results are shown in the Fig. 3
 147 and Fig. 4. It can be observed that Mamba fails to solve the inverse sequence matching task under
 148 nearly all configurations.

149 For the two-layer Transformer, results under different initialization schemes($\gamma = 0.5, 0.6, 0.7, 0.8, 0.9, 1.0$) are shown in the Fig. 5 and Fig. 6. It can be seen that even with
 150 a small number of layers, the Transformer outperforms Mamba.
 151




Figure 2: Loss of Mamba under fully symmetric setting. The horizontal axis represents the random seed, while the vertical axis corresponds to the number of layers in the Mamba model.

152 The results of the modified two-layer Mamba under different initialization schemes($\gamma =$
 153 $0.5, 0.6, 0.7, 0.8, 0.9, 1.0$) are shown in the Fig. 5 and Fig. 6. It can be observed that the modi-
 154 fied Mamba not only significantly outperforms standard Mamba, but also substantially surpasses the
 155 performance of the Transformer.




Figure 3: Accuracy of Mamba under different configurations on the inverse sequence matching task. The horizontal axis represents the number of layers in the Mamba model, while the vertical axis corresponds to the initialization scheme.




Figure 4: Loss of Mamba under different configurations on the inverse sequence matching task. The horizontal axis represents the number of layers in the Mamba model, while the vertical axis corresponds to the initialization scheme.




Figure 5: Accuracy of Mamba, Transformer, and modified Mamba on the inverse sequence matching task. Left: Mamba (2-layer), Center: Transformer (2-layer), Right: Modified Mamba (2-layer); all under varying initialization schemes.




Figure 6: Loss of Mamba, Transformer, and modified Mamba on the inverse sequence matching task. Left: Mamba (2-layer), Center: Transformer (2-layer), Right: Modified Mamba (2-layer); all under varying initialization schemes.